

Transfer Learning in Deep Reinforcement Learning

Author: **David Liang**
Advisor: **Professor Yu Hen Hu**

A thesis presented for
Undergraduate Honors in Computer Sciences



Department of Computer Sciences
University of Wisconsin-Madison
United States
5/10/2018

Abstract

Reinforcement learning has quickly risen in popularity because of its simple, intuitive nature and its powerful results. In this paper, we study a number of reinforcement learning algorithms, ranging from asynchronous q-learning to deep reinforcement learning. We focus on the improvements they provide over standard reinforcement learning algorithms, as well as the impact of initial starting conditions on the performance of a reinforcement learning agent.

Introduction

Reinforcement learning is a class of machine learning algorithms that are designed to allow agents provided with only the knowledge of the states it visits and the actions available to the agent to learn how to maximize its reward function, quite similar to the trial-and-error approach [3]. There are different techniques used for reinforcement learning, one of the most popular ones being “Q-learning” [14] where an agent develops a policy that chooses the action that is estimated to lead to the greatest total future rewards. Reinforcement learning has seen great recent success, particularly in “Playing Atari with Deep Reinforcement Learning” [7] and “Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm” [10] (AlphaZero) as it is a relatively simple yet extremely powerful algorithm, making it an interesting class of learning algorithms to study. Furthermore, the training of reinforcement learning agents is extremely slow, since the information it is provided is minimal, which means that there is a lot of room for improvement with reinforcement learning algorithms. Transfer learning on the other hand is a class of machine learning algorithms that seeks to transfer “knowledge” gained from solving one problem and applying it to another problem, so transfer learning can solve the problem of speed for reinforcement learning agents. In this paper, we discuss the impact of initial conditions with transfer learning on the convergence of reinforcement learning agents.

Motivation

In real life, we know that initial starting conditions matter. Consider a person who chooses to learn a sport: the athletic ability,

age, equipment, training, and instructor will all influence the time it takes for the person’s skill to peak. If it were at all possible, we would want to transfer the traits of high-performing athletes to the beginner to provide better chances at performing well. Based on this intuition, we want to experiment with transferring the models of trained reinforcement learning agents as the initial starting conditions of reinforcement learning algorithms and confirm that this hypothesis does indeed apply here too. That is, we want to show that given “better” initial conditions, an agent will likely achieve high performance faster than an agent with “worse” initial conditions. This is reasonable, and we can easily produce simple examples that illustrate the point. Consider an extreme example where an agent uses a neural network to model its policy, and all the weights in the network are initialized to zero. Then all the weights follow the same gradients, and the policy will likely perform poorly. Conversely, an agent with a policy model that has been trained for extremely long periods of time will likely be much closer to optimality: hence it will likely take much less time to converge. Intuitively, it makes sense that better initial conditions lead to optimal performance faster, and we wish to establish this for reinforcement learning, by means of a simple form of transfer learning.

Background

Reinforcement Learning

The reinforcement learning task is often formulated as a Markov decision process, a modeling framework useful for partially random, partially controlled environments, which is certainly the case in reinforcement learning where the environment may behave randomly, but the agent has control over its own actions. In the reinforcement learning task, a Markov decision pro-

cess consists of the following elements:

1. S_E : The set of states that the environment (with the agent in it) E can be in.
2. A_E : The set of possible actions that the agent can take in the environment.
3. $W_E : S_E \times A_E \rightarrow S_E$: The function that determines the resulting state given a starting state and an action.
4. $R_E : S_E \times S_E \rightarrow \mathbb{R}$: The function that gives the immediate reward for a state transition.

The agent constructs a policy $\pi_E : S_E \rightarrow A_E$ that maps a state in the state space to an available action that leads to the highest total immediate and future rewards. But because future rewards are not easily estimated in stochastic environments, a discount rate γ is introduced for future rewards so that given the choice between one immediate or one future reward (both equal in value), it will choose the immediate reward because of the stochastic environment. So we formulate a utility function $U_{\pi_E} : S_E \rightarrow \mathbb{R}$ that determines all rewards received by following the policy given a starting point s_0 :

$$U_{\pi_E}(s_0) = \sum_{t=0}^{\infty} \gamma^t R_E(s_t, W_E(s_t, \pi_E(s_t)))$$

Then the policy for our reinforcement learning agent can be defined as follows:

$$\pi_E(s_t) = \operatorname{argmax}_{a \in A_E} U_{\pi_E}(W_E(s_t, a))$$

Q-learning

Often times, an agent does not have access to W_E and in such cases, the agent's policy is said to be "model-free." The agent must, then, estimate the utility function by its internal Q-value function $Q_{\pi_E} : S_E \rightarrow \mathbb{R}$. A simple way to represent Q_{π_E} is to use a table in which each possible state and action pair is listed, and the estimated cumulative reward is the entry. To learn the optimal Q-value function which we denote by \hat{Q}_E , we use the Q-learning algorithm on our Q-value function Q_{π_E} . In one-step Q-learning, the algorithm takes one step at every training iteration t from state s_t , observes the reward received r_t and the new state s_{t+1} , and updates the policy as follows:

$$Q_{\pi_E}(s_t) = Q_{\pi_E}(s_t) + \alpha(r_t + \gamma Q_{\pi_E}(s_{t+1}) - Q_{\pi_E}(s_t))$$

with α as the learning rate, typically a real number between 0 and 1. This algorithm sets the target value to be the discounted sum of all the future rewards estimated by $\gamma Q_{\pi_E}(s_{t+1})$ added to the observed immediate reward r_t . The difference between the target value and output value is then weighted by the learning rate, and used to update the Q-function. To avoid settling for a non-optimal policy (premature convergence of policy), an exploration factor ϵ is introduced: ϵ is the probability that the agent will ignore its policy and execute a random action, to diversify its experiences and to avoid local minima in its policy. As time progresses, the exploration rate is decayed, so that the agent relies more (but not completely) on its policy. However, it is often the case that the exploration rate is not allowed to decay to 0 and is instead held at some fixed minimum exploration rate, to discourage the policy from sinking into a local minimum.

Q-networks

It becomes hard to maintain such a Q-table when the size of the state space increases: for example, consider an agent playing a video game, using the screen's pixel values as its state space. If the state is a $84 \times 84 \times 3$ array of 8 bit pixels and there are four actions available, the q-value table will hold $2^{84 \cdot 84 \cdot 3 + 2} \approx 10^{6351}$ entries! A popular solution to the problem of poorly scaling tables is the use of artificial neural networks in Q-learning termed Q-networks. Q-networks map states in the state space, represented by frames from the game, to q-values for each possible action. Q-networks learn to approximate \hat{Q} in a way intuitively similar to the update formula for the Q-table, by computing gradients for the network based on the output of the network (determined without knowing the next states) and target Q-values (determined using the next states) for the network. Q-networks are far more powerful than Q-tables because they can also approximate Q-values for states it has not yet seen and scales much better in terms of size. However, they come with the downside of being harder and slower to train.

Convolutional Neural Networks

One particular class of neural networks that has seen great success with images is the convolu-

tional neural network (CNN), used for the task of image classification [4] [17] [11] [12]. Part of its success is due to the fact that each layer can contain fewer weights, allowing for faster training. Each set of weights in a particular layer can be seen as a filter that is applied over the input image. A set of weights can then learn to detect certain features in each image of the input, and subsequent layers can learn mappings from these features to new features. A CNN also subsamples between layers to reduce the dimensionality of the image input, and is usually completed by a set of densely connected layers, the same as those found in the artificial neural network.

Deep Reinforcement Learning

With the above features, CNNs have become immensely popular, and lots of work has been done to apply CNNs to other fields. In particular, CNNs have been applied extremely successfully to reinforcement Q-learning, perhaps most notably in the deep Q-network (DQN) [7] [8]. This research has introduced a number of improvements. First, states are provided to the DQN agent as a sequence of frames, not just one. This provides temporal knowledge to the agent, which has proven to be beneficial. Then in order to combat the highly correlated data that comes from online learning, “experience replay” is introduced. This feature essentially stores a tuple of a state, the action taken from the state, the reward incurred for the action, and the resulting state from the action, termed an “experience.” In training, the experiences are randomly sampled and used for training the network. To improve the stability of the learning algorithm, every C updates, the network Q_{π_E} is cloned to obtain a target network \tilde{Q}_{π_E} , which is used to compute the Q-value targets. The gradients obtained are then applied to Q_{π_E} , not \tilde{Q}_{π_E} . Hence this less-frequently updated target network gives more stable gradients, leading to more stable learning.

Transfer Learning

As discussed before, it is well known that reinforcement learning can be slow, and in fact, one may end up running the same training algorithm for different environments, regardless of similarity in task, incurring an increase in training time linearly proportional with the

number of environments. As such, it would be advantageous if one model could be trained and have its knowledge transferred to the other models [13] [2] [15] [16]. This is the intuition behind transfer learning. Transfer learning gets quite complicated when it comes to determining what to transfer and how to transfer it. A simple example of transfer learning can be observed by training the agent to solve all of the subtasks that compose a single master task. This effectively allows the agent to complete the master task where it can be presented with any of the subtasks. In our particular case, we transfer the model of a pre-trained DQN agent to a beginning DQN agent and study its training with this initial condition.

Related work

Asynchronous Q-learning

An interesting improvement to Q-learning is asynchronous Q-learning (AQL). This technique involves one central, shared neural network. Then each asynchronous agent copies the shared network as its own individual network, learns on its own, and periodically shares its accumulated updates (gradients) with the shared neural network. Furthermore, each agent will periodically copy the shared neural network as its own individual neural network, making use of the learning that other agents have done. In effect, an AQL agent searches across multiple locations in the state space while sharing information with other agents, speeding up its learning process [6]. Another effect of AQL is that local minima are more easily avoided. This is due to the fact that there are multiple agents in different locations of the state space: an agent in a local minima can be pulled out when it copies in the shared neural network. Similarly, the shared neural network can be pulled out of a local minima when other agents share its updates with the shared neural network. Finally, AQL gives the improvement of CPU computation being more efficient than GPU computation, which is monetarily expensive [6]. As such, we see that AQL is a highly effective improvement to Q-learning.

Approach

We first describe the infrastructure available to us for our experiments. For high numbers of independent experiments, we use a distributed high-throughput computing resource through the Center for High Throughput Computing (CHTC) available at UW-Madison. As a result, we can run many independent experiments with lots of hardware resources, but we don't have GPU access, so training on CHTC is quite slow. For our asynchronous training experiment, we use OpenAI's "FrozenLake-v0" environment and consider the task to be solved when a win-rate of 78% is achieved, as defined by OpenAI. For our guaranteed convergence experiment, we tested using a custom maze environment with a state size of 4 and an action space size of 4. For our initial conditions experiment, we describe how we define an operational convergence criteria for our problem setting. We say that the agent's learning has stopped if the winning rate over the last 100 evaluations averages to a value greater than .78, the same stopping criteria for the environment FrozenLake.

Experiments

We conducted a number of experiments to study the state-of-the-art in reinforcement learning and to try to answer our initial questions. We present them below, describing our motivation and the results. Code for these experiments can be found in our [Github repository](#).

Asynchronous Training

We know from Mnih et al.[6] that asynchronous q-learning improves training speeds. Our hypothesis here was that an asynchronous q-network would provide faster training times (in terms of training epochs), proportional to the number of asynchronous agents used. Our hypothesis was affirmed by our experiment as seen in figure 1. It is worth noting, though, that our experiments were not run over a distributed computing system: instead, each agent was run on a single CPU, meaning that the total training times (in terms of pure time) were similar as shown in figure 2. We are confident that running the network over a large distributed system provide training times

inversely proportional to the number of agents used because of the inversely proportional change in iterations per agent. When running each agent on a separate core/cluster, these improvements in training time can be achieved. This experiment was conducted on a system with an Intel i7-7700k with 32 GB DDR4 3200 MHz SDRAM on a Samsung 960 EVO M.2 drive.

Guaranteed Convergence Given Infinite Time

We know from Even-Dar et. al[1] that using the action-elimination algorithm, our reinforcement learning agent will converge given infinite time. Our hypothesis was that the algorithm would work for a reinforcement learning agent in an environment with an extremely simple problem with an extremely small state space. We expected to see the algorithm converge given a couple months time. Unfortunately, the algorithm's progress exponentially diminished, and we never saw the convergence (or anything even close) after 2 months of running the algorithm on a high-throughput computing cluster. As such, we affirm that "infinite time" really does mean some enormous time quantity that is infeasible. We ran this experiment on a Google Cloud Compute Engine instance with 8 cores and 32 GB RAM.

Impact of Initial Conditions on Convergence

We hope to find that given "better" initial conditions, our DQN agent will converge faster. We provided these initial conditions as trained DQN models, saved after various periods of pre-training. We hypothesize that models that have had more pre-training will require less time to converge, while models that have had little pre-training. We first show the baseline performances of each initial condition in figure 3. Then we show training times until convergence starting from each initial condition in figure 4 (note that the experiments have not yet concluded so the figure is incomplete for the time being). Our hypothesis is affirmed through this experiment as we can see that indeed agents with more pre-training had faster times to convergence. The pre-training was done on a

Asynchronous Q-Learning Iteration Comparison

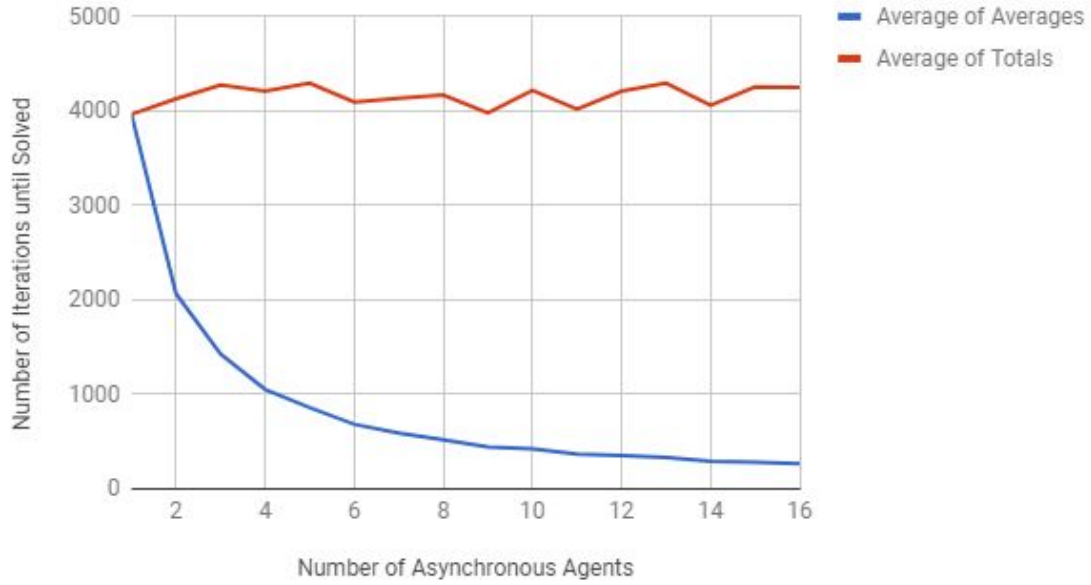


Figure 1: Comparison of average total number of iterations over all agents until the task was solved and the average of the number of iterations of each agent.

Asynchronous Q-Learning Time Comparison

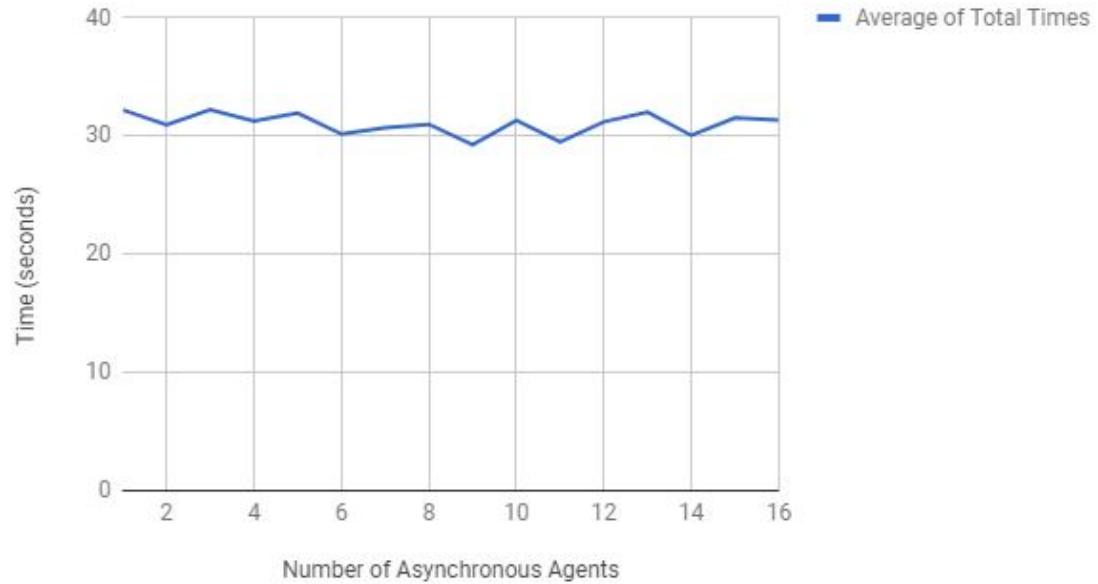


Figure 2: Comparison of average total time over all agents until the task was solved.

system with an Intel i7-7700k overclocked to 4.9 GHz with 32 GB DDR4 3200 MHz SDRAM on a Samsung 960 EVO M.2 drive. When testing each initial condition, we used CHTC. Each job was run on a system with 8 CPUs and 10 GB memory. Each job ran inside of a docker container.

Baseline Performances of DQN Agent

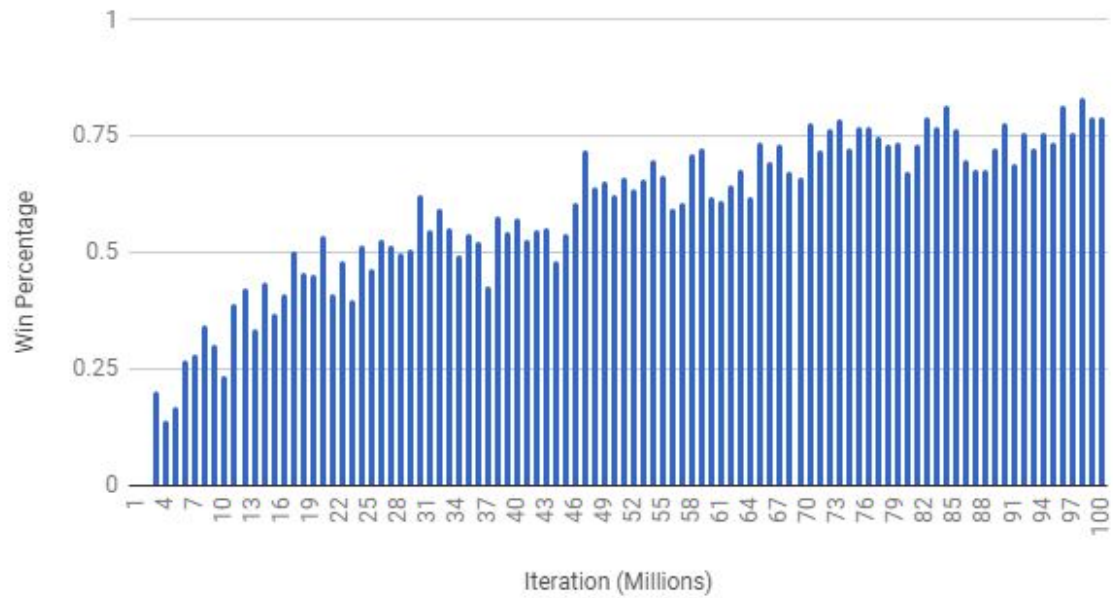


Figure 3: Baseline performances of the DQN agent. The x-axis shows the number of iterations that had passed when the agent was saved while the y-axis shows the winning percentage of the agent over 1000 games.




Figure 4: Comparison of time to convergence for different amounts of pretraining for DQN agents.

Conclusion

We have shown that initial conditions greatly impact the rate of convergence for reinforcement learning. As a result, transfer learning shows great potential for accelerating the convergence rate of reinforcement learning agents. Transfer learning has already seen great success in deep reinforcement learning, and we hope that this research is now further motivated.

Future Work

In the future, we would want to study adversarial learning in the reinforcement learning setting [5]. Intuitively, presenting challenges allows humans to learn better, and we believe that this translates to reinforcement learning agents as well. In fact, it has already been shown that this adds robustness [9]. Furthermore, adversarial learning is perfectly suited for two-player games like many of the Atari games. Hence our future work should include studies in adversarial learning in the reinforcement learning setting. We would also like to study learning models for multiple games and using transfer learning to apply these models to different reinforcement learning tasks.

References

- [1] Eyal Even-Dar, Shie Mannor, and Yishay Mansour. Action elimination and stopping conditions for reinforcement learning, 2006.
- [2] Meng Fang, Jie Yun, and Xingquan Zhu. Transfer learning across networks for collective classification, 2014.
- [3] Leslie Pack Kaelblinga, Michael L. Littman, and Andrew W. Moore. Reinforcement learning: A survey, 1996.
- [4] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. Imagenet classification with deep convolutional neural networks, 2012.
- [5] Yen-Chen Lin, Zhang-Wei Hong, Yuan-Hong Liao, Meng li Shih, Ming yu Liu, and Min Sun. Tactics of adversarial attack on deep reinforcement learning agents, 2017.
- [6] Volodymyr Mnih, Adrià Puigdomènech Badia, Mehdi Mirza, Alex Graves, Tim Harley, Timothy P. Lillicrap, David Silver, and Koray Kavukcuoglu. Asynchronous methods for reinforcement learning, 2016.

- [7] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin Riedmiller. Playing atari with deep reinforcement learning, 2013.
- [8] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dhharshan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. Human-level control through deep reinforcement learning, 2014.
- [9] Lerrel Pinto, James Davidson, Rahul Sukthankar, and Abhinav Gupta. Robust adversarial reinforcement learning, 2017.
- [10] David Silver, Thomas Hubert, Julian Schrittwieser, Ioannis Antonoglou, Matthew Lai, Arthur Guez, Marc Lan-
tot, Laurent Sifre, Dhharshan Kumaran, Thore Graepel, Timothy Lillicrap, Karen Simonyan, and Demis Hassabis. Mastering chess and shogi by self-play with a general reinforcement learning algorithm, 2017.
- [11] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition, 2014.
- [12] Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, and Andrew Rabinovich. Going deeper with convolutions, 2015.
- [13] Matthew E. Taylor and Peter Stone. Transfer learning for reinforcement learning domains: A survey, 2009.
- [14] Christopher John Cornish Hellaby Watkins. Learning from delayed rewards, 1989.
- [15] Ying Wei, Yu Zhang, and Qiang Yang. Learning to transfer, 2017.
- [16] Haiyan Yin and Sinno Jialin Pan. Knowledge transfer for deep reinforcement learning with hierarchical experience replay, 2017.
- [17] Matthew D. Zeiler and Rob Fergus. Visualizing and understanding convolutional networks, 2013.